

I can't directly export the information to your Google Docs, but I can provide the complete architecture and requirements set in a **copy-paste-ready Markdown format** that is structured for immediate use as a foundation for a Firebase/backend development project.

This document formalizes the software needed to manage the high volatility and "make-good" risk in your baseball club's advertising inventory.

Intelligent Ad Inventory & Risk Management System (AIMS) Requirements

1. System Overview and Goal

The primary goal of the Ad Inventory & Risk Management System (AIMS) is to transform the team's ad inventory from a static, volatile product into a **dynamic, data-driven financial asset**. This is achieved by using predictive modeling to quantify and proactively manage the risk of **viewership under-delivery (make-goods)** and **inventory shortage (game length)**.

2. Core Data Models (Relevant for Firestore/Firebase)

The system is powered by three main data models, which should be structured as top-level collections in a database like Firestore.

Model / Collection	Key Fields	Purpose
Games	game_id, date, opponent, home_team_win_prob, expected_duration_min, actual_duration_min, risk_index	Stores the schedule, predictive inputs, and the Risk Index score.
Contracts	contract_id,	The single source of truth

	advertiser_name, guaranteed_impressions, guaranteed_spots, base_cpm, total_spots_purchased, start_date, end_date	for all contractual obligations and pricing logic.
Ad_Delivery_Logs	log_id, game_id, spot_time, delivered_impressions, spot_type (Linear, Streaming, Virtual), make_good_liability_calc	Real-time record of what was actually shown to viewers, essential for reconciliation.
Inventory_Pool	inventory_date, spot_type, total_available_spots, reserved_make_good_spots, sellable_spots	Tracks all available inventory, ensuring the system knows exactly how much is left to sell and how much is reserved for liabilities.

3. Functional Modules & Requirements

The system must support three core modules with the following API-level capabilities:

A. AI Prediction Engine (Backend Services / Cloud Functions)

Requirement (R)	Description	Outcome
R3.1: Viewership Forecasting	Deploy a machine learning model to predict the ECR (Expected Commercial Rating) for every game 7, 3, and 1 day out, incorporating team performance, rivalry, and	Drives dynamic pricing for unsold inventory.

	day of week.	
R3.2: Game Duration Prediction	Deploy a model using historical and real-time data (pitcher, umpire, pitch clock) to forecast the final game length in minutes.	Accurately predicts the Total Available Inventory (spots) before the game starts.
R3.3: Risk Index Calculation	Automatically run a daily calculation to determine the Risk Index for all future games by comparing predicted inventory vs. contractual obligations.	Proactively flags games with a high probability of generating a make-good liability .

B. Dynamic Pricing & Sales Interface (API Endpoints)

Requirement (R)	Description	Outcome
R3.4: Dynamic Price Floor API	Expose an endpoint that the ad server queries in real-time to get the minimum price floor for floating ad spots, dynamically adjusted by the R3.1 forecast.	Maximizes revenue by ensuring high-demand spots are sold at a premium and low-demand spots are sold at an optimized discount.
R3.5: Inventory Allocation Service	Automatically pull inventory from the <code>Inventory_Pool</code> and mark a volume of spots as <code>reserved_make_good_spots</code> based on the R3.3 Risk Index.	Mitigates future risk by protecting inventory that will be needed to fulfill current shortfalls.
R3.6: Multi-Asset	Allow the sales team to	Creates diversified

Packaging	package linear broadcast spots with guaranteed digital/streaming ad inventory (e.g., pre-roll, virtual signage).	packages that reduce the dependency on the linear broadcast's volatility.
------------------	---	---

C. Make-Good and Reconciliation Module (Post-Game Automation)

Requirement (R)	Description	Outcome
R3.7: Automated Delivery Audit	Trigger a Cloud Function 3 hours after a game ends to ingest the Ad_Delivery_Logs and compare them against the Contracts guarantees.	Provides real-time, objective calculation of delivery performance for all advertisers.
R3.8: Make-Good Liability Service	Automatically calculate the financial liability (opportunity cost at base price) for the shortfall and convert it into a Make-Good Spot/Impression Credit .	Standardizes the cost of failure (e.g., the $\$130,000$ cost from our simulation) and ensures immediate liability recognition.
R3.9: Make-Good Fulfillment Report	Generate an executive report recommending specific future spots from the Make-Good Reserve Pool that should be given to the impacted advertiser to clear their credit.	Automates risk resolution and ensures the most cost-effective inventory is used for make-goods.

4. Technical Requirements & Firebase Focus

Component	Requirement	Firebase/GCP Service
Database	Scalable, real-time NoSQL database to store Games, Contracts, and Ad_Delivery_Logs data.	Cloud Firestore
Prediction Engine	Serverless environment for running Python/ML models for R3.1, R3.2, and R3.3 without managing infrastructure.	Cloud Functions or Cloud Run
Automation/APIs	Serverless functions triggered by events (game end, new contract) to run calculations (R3.7, R3.8).	Cloud Functions
Real-Time Data	Mechanism to ingest and synchronize high-velocity, low-latency game-state data.	Firebase Realtime Database or Google Pub/Sub
Hosting/UI	Simple, fast hosting for the internal sales and operations dashboard (the Risk Management Console).	Firebase Hosting